

NANYANG TECHNOLOGICAL UNIVERSITY

**NEXT-GEN GENERATIVE SEARCH ENGINE FOR CODE
BASE**

Timothy Lee Hongyi

College of Computing and Data Science

2026

NANYANG TECHNOLOGICAL UNIVERSITY

SC4079

NEXT-GEN GENERATIVE SEARCH ENGINE FOR CODE BASE

Submitted in Partial Fulfilment of the Requirements
for the Degree of Bachelor of Computing in Data Science and Artificial Intelligence
of the Nanyang Technological University

by

Timothy Lee Hongyi

College of Computing and Data Science

2026

Abstract

This report presents the design, implementation and evaluation of CodeOrient, an Artificial Intelligence (AI) Search tool, designed to accelerate developer onboarding. This application leverages Natural Language Processing (NLP), code graph visualisation and vector-based retrieval to help new developers understand unfamiliar codebases and reduce the time to first commit.

CodeOrient integrates semantic code search using embedding models, structural analysis through dependency graphs and generative user interfaces (UI) to provide context-aware feature cards. Through the use of Retrieval-Augmented Generation (RAG) with source grounding, the application eliminates hallucination, commonly seen in AI code assistants. Preliminary testing suggests that by externalising the mental model of a codebase through a unified graph-and-card interface, the system significantly reduces the cognitive load associated with onboarding and system discovery in large-scale repositories.

Acknowledgments

I would like to express my deepest gratitude to my supervisor, Professor Tan Chee Wei, for his unwavering guidance and patience throughout the duration of this Final Year Project. His expertise and insightful suggestions were key in shaping the direction of this project. I am very honoured to have had the opportunity to work under his mentorship.

Contents

Abstract	i
Acknowledgments	ii
Contents	iii
List of Tables	xi
List of Figures	xii
1 Introduction	1
1.1 Motivation and Problem Statement	1
1.2 Research Objectives and Goals	2
1.3 Key Contributions	2
1.4 Report Structure	3
2 Literature Review	4
2.1 Software Complexity Metrics and Code Quality Measurements	5
2.1.1 Cyclomatic Complexity as Foundation	5
2.1.2 Modern Complexity in Distributed Systems	6
2.2 Code Search and Natural Language Query Processing.....	6
2.2.1 Semantic Code Search Through Embeddings	6
2.2.2 Structural Code Search with Domain-Specific Languages	7
2.2.3 Code Embedding Models for Retrieval	7
2.3 AI Hallucination and Grounded Code Comprehension	8
2.3.1 The Hallucination Problem in Code-Generating LLMs	8
2.3.2 Citation-Grounded Code Comprehension	9

2.3.3	Retrieval-Augmented Generation	9
2.4	Code Visualisation and Dependency Analysis.....	9
2.4.1	Static Analysis and Dependency Graphs	9
2.4.2	Interactive Visualisation Tools	10
2.5	Generative User Interface (UI)	10
2.5.1	Generative UI as a Paradigm.....	10
2.5.2	Task-Driven Data Models for Adaptive Interfaces	10
2.6	Related Systems and Tools.....	11
2.7	Research Gaps and Motivation	11
2.7.1	Bridging Semantic and Structural Code Search.....	11
2.7.2	The "Black Box" of Generative UI in Software Engineering	11
2.8	Conclusion.....	12
3	System Design and Architecture	13
3.1	System Overview	14
3.1.1	Core Components	14
3.1.2	Data Flow Architecture.....	14
3.1.3	Microservices Structure	14
3.2	Frontend Architecture	14
3.2.1	Technology Stack	14
3.2.2	User Interface Design	14
3.2.3	Interactive Graph Visualization with React Flow	14
3.2.4	Generative UI Card Components	14
3.3	Backend Architecture	14
3.3.1	API Design.....	14
3.3.2	Code Search Engine	14
3.3.3	Dependency Graph Generation	14
3.3.4	LLM Integration and RAG Pipeline	14
3.4	Data Pipeline	14
3.4.1	Repository Ingestion	14
3.4.2	Code Parsing and AST Analysis	14

3.4.3	Embedding Generation	14
3.4.4	Dependency Graph Construction	14
3.4.5	Indexing and Storage	14
3.5	Natural Language Processing Pipeline	14
3.5.1	Query Understanding	14
3.5.2	Semantic Search using Embeddings	14
3.5.3	Structural Search using DSL Translation	14
3.5.4	Hybrid Retrieval Strategy	14
3.6	Citation-Grounded Generation	14
3.6.1	RAG Architecture	14
3.6.2	Source Code Grounding	14
3.6.3	Citation Mechanism	14
3.7	Technical Stack Justification	14
3.7.1	Frontend: Next.js	14
3.7.2	Backend: Next.js API Routes	14
3.7.3	LLM Integration: Vercel AI SDK and OpenAI	14
3.7.4	Code Search: GitHub Search API	14
3.7.5	Graph Visualization: React Flow	14
3.7.6	Database: PostgreSQL	14
3.7.7	Caching: Upstash Redis	14
3.7.8	Deployment: Vercel	14
4	Implementation Details	15
4.1	Development Methodology	16
4.1.1	Iterative Development Process	16
4.1.2	Version Control and Branching Strategy	16
4.2	Core Implementation Components	16
4.2.1	Search Module	16
4.2.2	Code Graph Analysis	16
4.2.3	Generative UI System	16
4.2.4	RAG Pipeline	16

4.3	API Endpoints and Data Models	16
4.4	Database Schema.....	16
4.5	Code Examples and Key Algorithms.....	16
4.5.1	Cyclomatic Complexity Calculation.....	16
4.5.2	Dependency Graph Construction	16
4.5.3	Embedding-Based Similarity Search	16
4.5.4	RAG Citation Grounding.....	16
4.6	Challenges and Solutions	16
4.6.1	Handling Large Codebases.....	16
4.6.2	Real-Time Graph Updates	16
4.6.3	Latency Optimization	16
4.6.4	Hallucination Prevention	16
4.6.5	Language and Framework Diversity	16
5	Evaluation Methodology	17
5.1	Evaluation Framework	18
5.1.1	Research Questions	18
5.1.2	Hypotheses.....	18
5.1.3	Evaluation Metrics	18
5.2	User Study Design	18
5.2.1	Participant Selection.....	18
5.2.2	Study Protocol	18
5.2.3	Task Design	18
5.2.4	Baseline Comparisons.....	18
5.3	Quantitative Metrics	18
5.3.1	Time-to-First-Commit.....	18
5.3.2	Time to Complete Onboarding Tasks.....	18
5.3.3	Code Search Precision and Recall	18
5.3.4	Graph Visualization Quality	18
5.3.5	Citation Accuracy and Hallucination Rate	18
5.4	Qualitative Metrics	18

5.4.1	User Satisfaction Survey	18
5.4.2	Usability Assessment.....	18
5.4.3	Cognitive Load Measurement.....	18
5.4.4	User Feedback and Interviews	18
5.5	Baseline Comparisons	18
5.5.1	Traditional Code Search (GitHub Search).....	18
5.5.2	Existing AI Code Assistants	18
5.5.3	Manual Code Navigation.....	18
5.6	Statistical Analysis	18
5.6.1	Hypothesis Testing	18
5.6.2	Effect Size Calculation	18
5.6.3	Confidence Intervals.....	18
6	Results and Analysis	19
6.1	Quantitative Results	20
6.1.1	Time-to-First-Commit Analysis	20
6.1.2	Search Performance Evaluation.....	20
6.1.3	Citation Accuracy Results.....	20
6.1.4	System Performance Metrics	20
6.2	Qualitative Results	20
6.2.1	User Satisfaction Findings	20
6.2.2	Usability Observations	20
6.2.3	Thematic Analysis of Feedback.....	20
6.3	Comparative Analysis with Baselines	20
6.3.1	Performance vs. GitHub Search	20
6.3.2	Performance vs. AI Code Assistants	20
6.3.3	Performance vs. Manual Navigation	20
6.4	Case Studies	20
6.4.1	Case Study 1: Understanding the Authentication Flow.....	20
6.4.2	Case Study 2: API Endpoint Discovery	20
6.4.3	Case Study 3: Error Handling Patterns.....	20

6.5	Data Visualisation and Interpretation	20
7	Discussion	21
7.1	Key Findings	22
7.1.1	Effectiveness of Citation Grounding.....	22
7.1.2	Impact of Code Graph Visualization	22
7.1.3	Generative UI Effectiveness.....	22
7.1.4	Hybrid Retrieval Strategy Benefits	22
7.2	Implications for Developer Onboarding.....	22
7.3	Limitations	22
7.3.1	Scale Limitations	22
7.3.2	Language and Framework Coverage.....	22
7.3.3	User Study Scope	22
7.3.4	Computational Resource Requirements	22
7.4	Comparison with Related Work	22
7.4.1	How We Advance Beyond Existing Solutions	22
7.4.2	Unique Contributions	22
7.5	Threats to Validity.....	22
7.5.1	Internal Validity	22
7.5.2	External Validity.....	22
7.5.3	Construct Validity	22
7.6	Practical Implications for Industry	22
8	Conclusion	23
8.1	Summary of Contributions	23
8.2	Key Takeaways	23
8.3	Future Work	23
8.3.1	Scalability Improvements	23
8.3.2	Multi-Language Support	23
8.3.3	Advanced Generative UI	23
8.3.4	Integration with Developer Tools	23
8.3.5	Real-Time Collaborative Features	23

8.4	Recommendations for Implementation.....	23
8.5	Final Remarks	23
A	User Study Materials	24
A.1	Participant Consent Form.....	24
A.2	Study Instructions	24
A.3	Task Descriptions	24
A.4	Survey Questionnaire	24
A.5	Interview Questions	24
B	Detailed Results Tables	25
B.1	Quantitative Results Tables	25
B.2	User Feedback Summary	25
B.3	Performance Metrics by Repository	25
C	Code Snippets and Implementation Details	26
C.1	Key Algorithm Implementations	26
C.2	API Endpoint Specifications.....	26
C.3	Database Schema SQL.....	26
C.4	Frontend Component Code	26
C.5	RAG Pipeline Code	26
D	System Architecture Diagrams	27
D.1	High-Level System Architecture	27
D.2	Data Flow Diagrams	27
D.3	Component Interaction Diagram	27
D.4	Deployment Architecture	27
E	Additional Evaluation Data	28
E.1	Raw User Study Data	28
E.2	Search Performance Analysis.....	28
E.3	Citation Accuracy Detailed Results	28
E.4	Latency Distributions	28

F Related Work Comparison Table	29
F.1 Feature Comparison Matrix	29
F.2 Performance Comparison	29

List of Tables

List of Figures

2.1	<i>Control flow graph of a simple if-else statement.</i>	5
2.2	<i>Semantic code search in a shared embedding space for retrieval [3].</i>	7
2.3	<i>Exploiting LLM hallucinations through slopsquatting.</i>	8
2.4	<i>Retrieval-Augmented Generation to reduce hallucinations [7].</i>	9
2.5	<i>Generative UI for a Room Rug Visualiser [11].</i>	12

Chapter 1

Introduction

Software Engineering is undergoing a fundamental shift with the rise in Large Language Models (LLMs), Agentic Artificial Intelligence (AI), and generative user interfaces (GenUI). As the complexity of current software architectures (distributed systems) and codebases grow, the challenge of developers learning a new repository has become a significant bottleneck for engineering productivity. I propose CodeOrient, an autonomous AI-driven developer onboarding platform that leverages Retrieval-Augmented Generation (RAG) with dynamic graph visualisation. Deployed as an interactive onboarding assistant, CodeOrient combines the reasoning capabilities of LLMs with the structural insights of code graphs to transform how developers understand and navigate unfamiliar codebases.

1.1 Motivation and Problem Statement

The rapid advancement of AI-assisted coding tools, such as GitHub Copilot and Cursor, has prioritised code generation over code comprehension. However, for a developer joining a new organisation, the primary hurdle is not writing new code, but understanding the existing codebases tied within complex architectures. Current onboarding processes face three critical issues:

1. The Hallucination Problem: LLMs often generate plausible-sounding but incorrect explanations of code structure, leading to confusion and mistrust.

2. The Context Window Constraint: Large codebases exceed the input limits of LLMs, resulting in fragmented or incomplete answers.
3. Cognitive Overload in Navigation: Static documentation fails to capture the dynamic relationships within code, forcing developers to mentally map text-based

1.2 Research Objectives and Goals

The primary goal of this research is to create a citation-grounded, visually adaptive system that reduces the time required for a new developer to familiarise with the codebases. The specific objectives are:

1. Develop a Retrieval-Augmented Generation (RAG) Framework: Moving beyond keyword/fuzzy search to a hybrid system that queries semantic embeddings to capture user intent.
2. Implement a Generative UI for Code Visualisation: Leveraging libraries like React Flow to dynamically render interactive graphs that adapt to the user's specific natural language intent.
3. Minimise AI Hallucinations through Citation Grounding: Ensuring every architectural claim made by the LLM is backed by code snippets or graph nodes from the actual repository.
4. Evaluate the Impact of Visual Context on Onboarding: To explore how dynamic graph visualisations affect developer comprehension and cognitive load compared to traditional text-based documentation.

1.3 Key Contributions

CodeOrient introduces several innovations that distinguish it from conventional AI coding assistants:

- Dynamic Graph Synthesis: Instead of text-based answers, Generative UI is used to create custom visualisations for each query. If a developer asks about “authentication flow,” the UI creates a graph focused strictly on those related modules, rather than a cluttered, static diagram.
- Citation-Grounded RAG Pipeline: By integrating citation mechanisms into the RAG framework, CodeOrient ensures that all LLM outputs are verifiable against actual code segments, significantly reducing misinformation.
- Agentic LLM Integration: CodeOrient utilises agentic LLMs capable of autonomously deciding when to query the vector database, generate visualisations, or seek clarifications. This autonomy allows for a more fluid interaction between the developer and the system.

1.4 Report Structure

This report is organised to guide the reader through the theoretical foundations and technical implementation of CodeOrient:

Chapter 2: Literature Review explores the intersection of software complexity metrics, semantic code search, and the emerging technologies of Generative UI.

Chapter 3: System Design and Architecture details the technical stack, including the integration of Agentic LLMs with vector databases and React Flow-based frontend.

Chapter 4: Implementation Details discusses the challenges of grounding AI outputs in the source code and the development of the RAG pipeline.

Chapter 5 & 6: Evaluation and Results will analyse the tool’s effectiveness in reducing cognitive load and improving search precision.

Chapter 7: Discussion discusses the implications of the findings, limitations of the current approach, and potential avenues for future research.

Chapter 8: Conclusion summarises the contributions of this research and reflects on the broader impact of AI-driven developer onboarding tools.

Chapter 2

Literature Review

Understanding unfamiliar and large codebases poses significant challenges for software developers, especially those new to a project or organisation. Different organisations have various ways to onboard new developers, yet these approaches to codebase exploration remains inefficient. As AI adoption increases, the problem of accurate code comprehension becomes paramount. Often, AI-assisted tools hallucinate information that is not grounded in actual source code. Furthermore, large codebases make it harder for AI to comprehend due to its limited context window.

This literature review examines five interrelated research domains critical to the proposed AI Search Tool - CodeOrient:

1. Software Complexity Metrics and Code Quality Measurements
2. Semantic Code Search and Natural Language Query Processing
3. AI Hallucination and Grounded Code Comprehension
4. Code Visualisation and Dependency Analysis
5. Generative User Interfaces

The combination of these areas provides the theoretical foundation for building an AI application that reduces developer onboarding time while maintaining citation accuracy and reliability.

2.1 Software Complexity Metrics and Code Quality Measurements

2.1.1 Cyclomatic Complexity as Foundation

Thomas J.M McCabe introduced cyclomatic complexity, a quantitative measure of program complexity based on control graph flow analysis [1]. His work established that cyclomatic complexity directly correlates with code maintainability and testing. The formula $M = E - N + 2P$, where E represents edges, N represents nodes, and P represents the number of connected components in a control flow graph, represents the complexity as the number of linearly independent paths through a program's source code [1]. As shown in Fig.1, a simple control flow graph of a function below yields a complexity of 2, where $P = 1$.

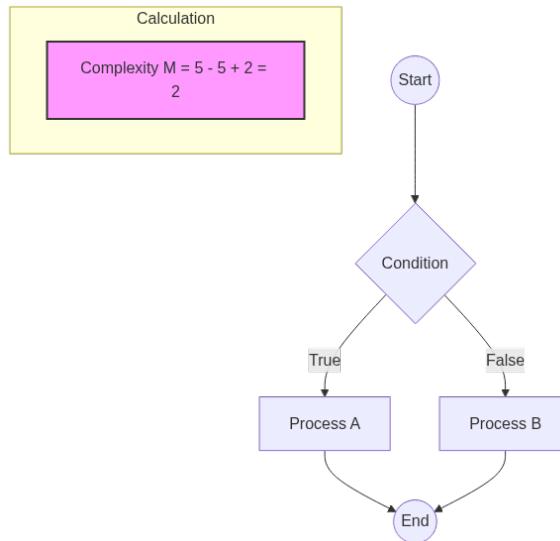


Figure 2.1: *Control flow graph of a simple if-else statement.*

McCabe's complexity measure allows developers to identify highly complex functions and recognise problematic code sections that require refactoring. When implementing the code graph visualisation feature in a code search application, cyclomatic complexity serves as one signal among many to highlight high-risk or critical code sections.

2.1.2 Modern Complexity in Distributed Systems

Kafura's recent reflection on McCabe's work in 2025 acknowledges that cyclomatic complexity has proven durable for the last 50 years. However, modern software architectures, such as distributed systems and microservices, require additional metrics beyond control flow analysis [2]. This observation directly motivates the integration of code graph visualisation in modern development tools, as visual representation of information flow across procedures becomes just as important as understanding control flow within individual functions.

2.2 Code Search and Natural Language Query Processing

2.2.1 Semantic Code Search Through Embeddings

With the recent advancements in neural networks and embedding models, the field of semantic code search has evolved significantly from traditional keyword-based search. Cambronero et al. demonstrated that the use of neural embeddings could bridge the semantic gap between natural language queries and code snippets [3]. By transforming both code and queries into a shared vector space, and code snippets relevant to the query can be retrieved by calculating the cosine similarity between their embedding vectors, given by the formula:

$$\text{cosine_similarity}(\vec{a}, \vec{b}) = \frac{\vec{a} \cdot \vec{b}}{\|\vec{a}\| \|\vec{b}\|}$$

where \vec{a} and \vec{b} are the embedding vectors of the query and code snippet, respectively.

In the figure below, developers can express their intent through natural language, which is then mapped into the same embedding space as code snippets. This technique supplements existing searches, such as fuzzy search or keyword-based approaches that often miss semantically relevant results. This motivates the use of natural language search in modern code repositories like GitHub.

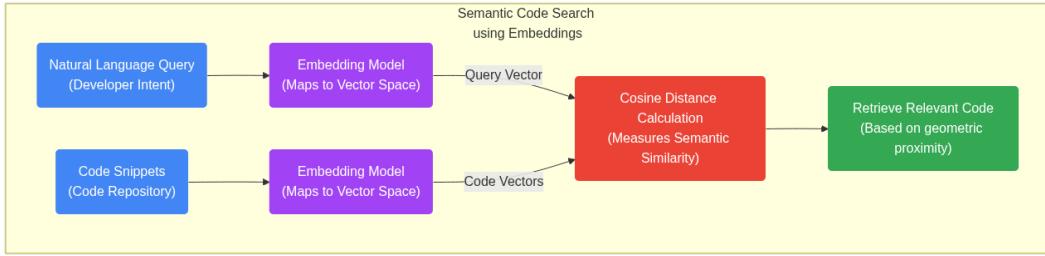


Figure 2.2: *Semantic code search in a shared embedding space for retrieval [3]*.

2.2.2 Structural Code Search with Domain-Specific Languages

Recent research by Limpanukorn et al. (2025) introduces structural code search by translating natural language queries into Domain-Specific Language (DSL) queries for more precise code retrieval [4]. This approach leverages the architectural information embedded in a codebase rather than just its textual meaning. This method achieved a precision score of 55-70% and outperformed semantic search baselines by up to 57% on F1 scores [4].

This research presents a critical missing link in developer tools, such as tracing an "authentication flow", where the relationships between modules are more informative than the functions' names themselves. The findings provide a strong validation for the code graph visualisation tool, providing syntactic and architectural information for developers to understand complex problem spaces.

2.2.3 Code Embedding Models for Retrieval

In 2025, Qodo introduced specialised code embedding models (Qodo Embed-1) that achieved state-of-the-art performance in Codebase Understanding Gartner® 2025, with a product score of 3.72/5. Their approach directly encodes code semantics without an intermediate language description step [5]. By avoiding the overhead of the intermediate step, Qodo Embed-1 has proved to be computationally efficient while maintaining high retrieval accuracy.

2.3 AI Hallucination and Grounded Code Comprehension

2.3.1 The Hallucination Problem in Code-Generating LLMs

With the rise in AI-assisted code generation tools such as GitHub Copilot and ChatGPT, the problem of hallucination becomes increasingly critical. Hallucination refers to generated information that seems plausible but is fabricated or factually incorrect. Spracklen et al.'s (2025) research revealed that package hallucinations are a systemic issue across state-of-the-art code-generating models [6]. They analysed over 576,000 generated code samples across 16 Large Language Models (LLMs) and found that the LLMs consistently hallucinate package names. More critically, they regenerate the same false package name across 43% of repeated prompts [6].

This poses a critical issue in agentic workflow as the hallucination is exploitable via "slopsquatting", creating risks in the software supply chain [6]. The figure below showcases how an attacker can exploit hallucinations from LLMs. This research highlights the urgent need for citation-grounded code comprehension systems that can verify all claims against actual source code.

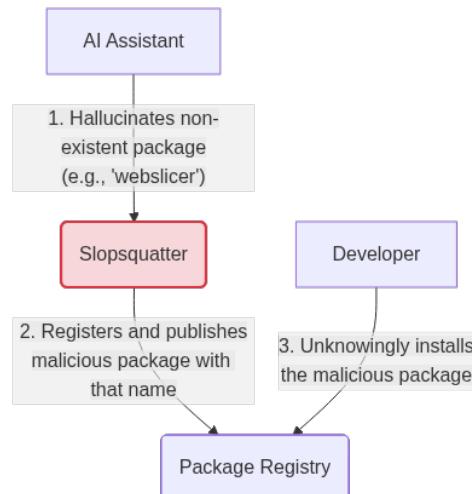


Figure 2.3: Exploiting LLM hallucinations through slopsquatting.

2.3.2 Citation-Grounded Code Comprehension

Arafat et al.’s recent work on citation-grounded code comprehension directly addresses the hallucination problem [7]. They conclude that code comprehension systems must ground all claims in verifiable source code citations. Their proposed hybrid retrieval system with Neo4j graph database to provide import relationships, achieved a 92% citation accuracy with zero hallucinations. Moreover, the graph component discovered richer cross-file relationships that purely text-based retrieval missed, 62% of architectural queries [7].

2.3.3 Retrieval-Augmented Generation

The broader principle emerging from hallucination research is Retrieval-Augmented Generation (RAG). As it is not possible to train new information into LLMs at scale, RAG provides a mechanism to ground LLM outputs with real-time data via a retriever [8]. The use of a retriever reduced hallucination rates across all categories from a baseline high of 21% to below 7.5% [8]. In the context of code comprehension, RAG refers to a retriever that fetches relevant code snippets, which are then passed to an LLM as a context to generate grounded explanations with source citations.

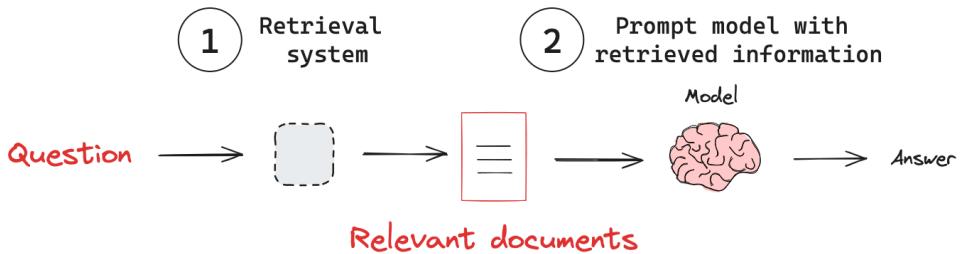


Figure 2.4: *Retrieval-Augmented Generation to reduce hallucinations* [7].

2.4 Code Visualisation and Dependency Analysis

2.4.1 Static Analysis and Dependency Graphs

To understand the real structure of a codebase, dependency graphs are essential. Using entities as nodes and relationships as edges, code graphs provide a visual representation

of how different components interact with each other [9]. Without this visualisation, it is inherently difficult for developers to grasp the complex relationships present in modern software architectures like microservices. In particular, for CodeOrient, visualising dependencies helps developers trace data flows and control flows across modules, which helps them understand unfamiliar codebases faster.

2.4.2 Interactive Visualisation Tools

React Flow is a popular open-source library for building interactive, node-based User Interfaces (UIs) in React applications [10]. This interactivity is crucial for developers to understand the linkage between nodes and edges and how they interact. Rather than traditional static diagrams, developers can explore relationships dynamically by hovering over nodes and their related code snippets.

2.5 Generative User Interface (UI)

2.5.1 Generative UI as a Paradigm

Google’s recent research on Generative UI introduces a paradigm shift in interface design [11]. Traditionally, LLMs can only generate text-based outputs, but Generative UI extends this capability by generating dynamic user interfaces at runtime based on user prompts. This strengthens the capability of LLMs by giving them the tools to generate dynamic code graphs for visualisation, with their structure and content optimised for each specific query.

2.5.2 Task-Driven Data Models for Adaptive Interfaces

Recent research by Cao et al. showcases the use of task-driven data models as the foundation for Generative UI [12]. By allowing LLMs to generate a data model representing the core task, then mapping it to UI specifications, the resulting interfaces were more adaptive and aligned with user intent than direct UI code generation. This also ensures that the generated UI is grounded in actual data structures rather than arbitrary code snippets.

2.6 Related Systems and Tools

Based on current semantic search tools like RAG [8] and interactive graph visualisation tools [10], there is a clear gap in integrating these capabilities into a unified developer onboarding experience. This addresses the onboarding challenges faced by developers when exploring unfamiliar codebases. The integration of LLMs, semantic search, code graph visualisation, and generative UI into a single application allows developers to search for relevant code snippets, understand their relationships via dynamically generated graphs, while ensuring trust through citation grounding [7].

2.7 Research Gaps and Motivation

While the individual domains of semantic search, graph analysis, and Generative UI have matured significantly, their integration into a developer tool reveals two critical research gaps that CodeOrient aims to address:

2.7.1 Bridging Semantic and Structural Code Search

Current tools typically favour either semantic search (finding code that looks right) or structural analysis (finding code that is connected). As noted by Limpanukorn et al. (2025), structural search outperforms semantic baselines, yet most AI tools like GitHub Copilot still relies primarily on text-based RAG. There is a lack of research into how Generative UI can bridge this gap by dynamically synthesising a visual graph that represents both the user’s natural language intent and the codebase’s physical architecture.

2.7.2 The ”Black Box” of Generative UI in Software Engineering

Research by Leviathan et al. [11] and Cao et al. [12] establishes the framework for task-driven UIs. However, these studies focus on general tasks such as education or shopping, as shown below. In the high-stakes domain of software engineering, it is unknown how a constantly changing, generative interface affects a developer’s productivity.

CodeOrient provides an experimental platform to observe whether generative graphs reduce cognitive load during onboarding or if the lack of visual consistency hinders comprehension.

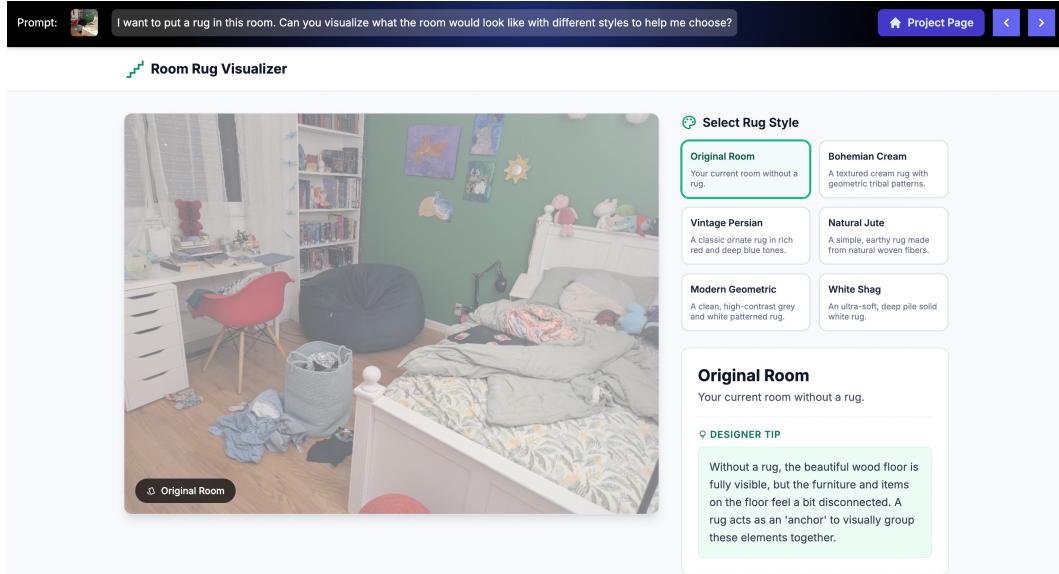


Figure 2.5: *Generative UI for a Room Rug Visualiser [11]*.

2.8 Conclusion

Improving developer onboarding requires addressing code comprehension at multiple levels. Recent research in cyclomatic complexity [1], semantic code search [3], citation-grounded AI [7], and Generative UI [12] provides a comprehensive foundation. By integrating these approaches together, CodeOrient has the potential to accelerate developer onboarding while being reliable and accurate.

Chapter 3

System Design and Architecture

3.1 System Overview

3.1.1 Core Components

3.1.2 Data Flow Architecture

3.1.3 Microservices Structure

3.2 Frontend Architecture

3.2.1 Technology Stack

3.2.2 User Interface Design

3.2.3 Interactive Graph Visualization with React Flow

3.2.4 Generative UI Card Components

3.3 Backend Architecture

3.3.1 API Design

3.3.2 Code Search Engine

3.3.3 Dependency Graph Generation¹⁴

3.3.4 LLM Integration and RAG Pipeline

Chapter 4

Implementation Details

4.1 Development Methodology

4.1.1 Iterative Development Process

4.1.2 Version Control and Branching Strategy

4.2 Core Implementation Components

4.2.1 Search Module

BM25 Sparse Retrieval

Dense Embedding Retrieval

Hybrid Search Integration

4.2.2 Code Graph Analysis

Dependency Extraction

Graph Construction Algorithm

Graph Traversal and Ranking

4.2.3 Generative UI System

Query-to-Task Model Translation 16

Task Model to UI Specification

Chapter 5

Evaluation Methodology

5.1 Evaluation Framework

5.1.1 Research Questions

5.1.2 Hypotheses

5.1.3 Evaluation Metrics

5.2 User Study Design

5.2.1 Participant Selection

5.2.2 Study Protocol

5.2.3 Task Design

5.2.4 Baseline Comparisons

5.3 Quantitative Metrics

5.3.1 Time-to-First-Commit

5.3.2 Time to Complete Onboarding Tasks

5.3.3 Code Search Precision and Recall¹⁸

5.3.4 Graph Visualization Quality

Chapter 6

Results and Analysis

6.1 Quantitative Results

6.1.1 Time-to-First-Commit Analysis

6.1.2 Search Performance Evaluation

6.1.3 Citation Accuracy Results

6.1.4 System Performance Metrics

Query Latency

Throughput

Resource Utilisation

6.2 Qualitative Results

6.2.1 User Satisfaction Findings

6.2.2 Usability Observations

6.2.3 Thematic Analysis of Feedback

6.3 Comparative Analysis with Baselines₂₀

6.3.1 Performance vs. GitHub Search

Chapter 7

Discussion

7.1 Key Findings

7.1.1 Effectiveness of Citation Grounding

7.1.2 Impact of Code Graph Visualization

7.1.3 Generative UI Effectiveness

7.1.4 Hybrid Retrieval Strategy Benefits

7.2 Implications for Developer Onboarding

7.3 Limitations

7.3.1 Scale Limitations

7.3.2 Language and Framework Coverage

7.3.3 User Study Scope

7.3.4 Computational Resource Requirements

7.4 Comparison with Related Work

7.4.1 How We Advance Beyond Existing Solutions

7.4.2 Unique Contributions

Chapter 8

Conclusion

8.1 Summary of Contributions

8.2 Key Takeaways

8.3 Future Work

8.3.1 Scalability Improvements

8.3.2 Multi-Language Support

8.3.3 Advanced Generative UI

8.3.4 Integration with Developer Tools

8.3.5 Real-Time Collaborative Features

8.4 Recommendations for Implementation

8.5 Final Remarks

Appendix A

User Study Materials

A.1 Participant Consent Form

A.2 Study Instructions

A.3 Task Descriptions

A.4 Survey Questionnaire

A.5 Interview Questions

Appendix B

Detailed Results Tables

B.1 Quantitative Results Tables

B.2 User Feedback Summary

B.3 Performance Metrics by Repository

Appendix C

Code Snippets and Implementation Details

C.1 Key Algorithm Implementations

C.2 API Endpoint Specifications

C.3 Database Schema SQL

C.4 Frontend Component Code

C.5 RAG Pipeline Code

Appendix D

System Architecture Diagrams

D.1 High-Level System Architecture

D.2 Data Flow Diagrams

D.3 Component Interaction Diagram

D.4 Deployment Architecture

Appendix E

Additional Evaluation Data

E.1 Raw User Study Data

E.2 Search Performance Analysis

E.3 Citation Accuracy Detailed Results

E.4 Latency Distributions

Appendix F

Related Work Comparison Table

F.1 Feature Comparison Matrix

F.2 Performance Comparison

Bibliography

- [1] T.J. McCabe. “A Complexity Measure”. In: *IEEE Transactions on Software Engineering* SE-2.4 (1976), pp. 308–320. doi: [10.1109/TSE.1976.233837](https://doi.org/10.1109/TSE.1976.233837).
- [2] Dennis Kafura. “Reflections on McCabe’s Cyclomatic Complexity”. In: *IEEE Transactions on Software Engineering* 51.3 (2025), pp. 700–705. doi: [10.1109/TSE.2025.3534580](https://doi.org/10.1109/TSE.2025.3534580).
- [3] José Cambronero et al. “When Deep Learning Met Code Search”. In: *CoRR* abs/1905.03813 (2019). arXiv: [1905.03813](https://arxiv.org/abs/1905.03813). URL: <http://arxiv.org/abs/1905.03813>.
- [4] Ben Limpanukorn et al. *Structural Code Search using Natural Language Queries*. 2025. arXiv: [2507.02107](https://arxiv.org/abs/2507.02107) [cs.SE]. URL: <https://arxiv.org/abs/2507.02107>.
- [5] Sheffer Tai. *State-of-the-Art Code Retrieval with Efficient Embeddings*. 2025. URL: <https://www.qodo.ai/blog/qodo-embed-1-code-embedding-code-retrieval/>.
- [6] Joseph Spracklen et al. *We Have a Package for You! A Comprehensive Analysis of Package Hallucinations by Code Generating LLMs*. 2025. arXiv: [2406.10279](https://arxiv.org/abs/2406.10279) [cs.SE]. URL: <https://arxiv.org/abs/2406.10279>.
- [7] Jahidul Arafat. *Citation-Grounded Code Comprehension: Preventing LLM Hallucination Through Hybrid Retrieval and Graph-Augmented Context*. 2025. arXiv: [2512.12117](https://arxiv.org/abs/2512.12117) [cs.SE]. URL: <https://arxiv.org/abs/2512.12117>.
- [8] Orlando Ayala and Patrice Bechar. “Reducing hallucination in structured outputs via Retrieval-Augmented Generation”. In: *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies* (NAACL-HLT 2024), pp. 1–11. 2024. doi: [10.1162/nacl_a_02007](https://doi.org/10.1162/nacl_a_02007).

- Linguistics: Human Language Technologies (Volume 6: Industry Track)*. Association for Computational Linguistics, 2024, pp. 228–238. doi: 10.18653/v1/2024.nacl-industry.19. URL: <http://dx.doi.org/10.18653/v1/2024.nacl-industry.19>.
- [9] Jit. *How to Use a Dependency Graph to Analyze Dependencies*. 2025. URL: <https://www.jit.io/resources/app-security/how-to-use-a-dependency-graph-to-analyze-dependencies>.
 - [10] Nikola Jovanov et al. “A visual approach to project management using react flow”. In: Jan. 2025, pp. 306–312. doi: 10.5937/IIZS25306J.
 - [11] Y. Leviathan et al. *Generative UI: LLMs are Effective UI Generators*. 2025. URL: <https://research.google/blog/generative-ui-a-rich-custom-visual-interactive-user-experience-for-any-prompt/>.
 - [12] Yining Cao, Peiling Jiang, and Haijun Xia. *Generative and Malleable User Interfaces with Generative and Evolving Task-Driven Data Model*. 2025. arXiv: 2503.04084 [cs.HC]. URL: <https://arxiv.org/abs/2503.04084>.